



**Department of Electrical & Computer Engineering**  
**Intelligent Systems Lab ENCS5141**

---

**Case Study #2:**

---

**Prepared by:**

**Names:** Abdelrhman Abed  
**ID:**1193191

**Instructor:**

Dr. Aziz Qaroush

**Sec: 2**

**Date: 8/1/2025**

## Abstract

This case study addresses the challenge of subjectivity detection in text, utilizing image-based character sequences for sentence reconstruction and natural language processing (NLP) techniques for classification. In Task 1, a Convolutional Neural Network (CNN) is trained on character image data to extract text from images, with additional evaluation using a pre-trained model for character recognition. The extracted sentences undergo preprocessing to prepare them for subsequent analysis. Task 2 focuses on the NLP task of classifying sentences as subjective or objective. A Long Short-Term Memory (LSTM) model is implemented and trained, and transfer learning techniques using pre-trained transformer models are applied to enhance performance. Both tasks are thoroughly evaluated using standard performance metrics, including accuracy, precision, recall, and F1-score. Hyperparameter tuning is conducted to optimize the models. The report analyzes the results, highlighting the strengths and limitations of the proposed methodologies for subjectivity detection.

## Table of Contents

<b>Abstract.....</b>	<b>II</b>
<b>1.Introduction.....</b>	<b>5</b>
<b>2. Procedure and Discussion .....</b>	<b>6</b>
<b>2.1 Task 1: Subjectivity Classification using CNN and Pre-Trained ResNet18 .....</b>	<b>6</b>
<b>2.2 Task 2: Subjectivity Classification using LSTM and Transfer Learning.....</b>	<b>8</b>
<b>3.3 Discussion of Results.....</b>	<b>10</b>
<b>3.conclusion.....</b>	<b>12</b>

## Table of Figures

FIGURE 1: DATA FILES USED IN TASK 1 .....	6
FIGURE 2: CNN MODEL RESULT .....	7
FIGURE 3: RESNET18 MODEL RESULT .....	7
FIGURE 4: DATA FILES USED IN TASK 2 .....	8
FIGURE 5: LSTM MODEL RESULT .....	8
FIGURE 6: TRANSFER LEARNING RESULT .....	9
FIGURE 7: HYPERPARAMETER TUNING RESULT .....	10
FIGURE 8: TRAINING LOSS COMPARISON .....	10

## Table of Tables

TABLE 1: PERFORMANCE COMPARISON OF DIFFERENT MODELS FOR CHARACTER RECOGNITION TASK .....	10
--	----

# 1.Introduction

Deep learning has significantly impacted the fields of image and text processing, enabling the development of sophisticated models for complex tasks. This project focuses on two tasks: image classification and subjectivity classification, leveraging powerful neural network architectures like Convolutional Neural Networks (CNNs) for image data and Long Short-Term Memory (LSTM) networks for text data.

In the first task, the goal is to classify handwritten character images using a custom CNN model. The dataset consists of images provided in files such as **characters.csv** (training data samples) and **characters-test.csv** (test data samples), with label mapping information from **mappings.txt** for encoding. The data preprocessing steps include loading and cleaning data (handling missing values and removing non-alphanumeric characters), label encoding, feature transformation into character-level and word-level vectors, and splitting the training dataset into training and validation subsets. CNNs are well-suited for this task due to their ability to learn spatial hierarchies of features. Additionally, a pre-trained model, such as ResNet or VGG, will be fine-tuned for character recognition, allowing for insights into the benefits of transfer learning.

In the second task, the objective is to classify text as subjective (opinion-based) or objective (factual). This involves implementing an LSTM model for sequential input with an architecture that includes an input layer for sequence embedding (word-based or character-based), LSTM layers for long-range dependency learning, and a dense layer with softmax activation for binary classification. Transfer learning will be applied by using a pre-trained embedding model, such as GloVe, to initialize the embedding layer and fine-tuning it on the training data for subjectivity classification.

The performance of the models will be evaluated using metrics such as:

- **Accuracy:** Measures overall correctness.
- **Precision:** Identifies the proportion of predicted subjective labels that are correct.
- **Recall:** Determines how many of the actual subjective samples were correctly predicted.
- **F1-score:** Balances precision and recall for a more comprehensive performance assessment.

This case study aims to explore deep learning techniques for both image and text classification, demonstrating the use of CNNs for image recognition and LSTMs for natural language processing. Transfer learning will be employed in both tasks to enhance model performance, showcasing how pre-trained models can be adapted for specific classification challenges. The outcome will provide valuable insights into the power of deep learning in solving real-world problems across multiple domains.

## 2. Procedure and Discussion

### 2.1 Task 1: Subjectivity Classification using CNN and Pre-Trained ResNet18

#### Data Files Used:

characters.csv (Training data with 112,799 samples and 785 features, where each row represents a text snippet in encoded format), characters-test.csv (Test data with 18,799 samples and 785 features, formatted similarly to the training data for model evaluation), mappings.txt (Mapping information for label encoding, providing the correspondence between numeric values and their respective class labels for subjectivity and objectivity).

```
Train data shape: (112799, 785)
Test data shape: (18799, 785)
Dropping the first column of test data.
X_train shape: (90239, 784), X_val shape: (22560, 784), test_features shape: (18799, 784)
Using device: cuda
```

*Figure 1: data files used in task 1*

#### Objective:

Train a CNN model for subjectivity classification by transforming text snippets into a structured input format and evaluating its performance across different metrics. Additionally, a pre-trained ResNet18 model is used to leverage deep feature extraction, demonstrating the effectiveness of transfer learning in improving classification accuracy and generalization on unseen data. The goal is to compare performance, training efficiency, and generalization capabilities between these models.

#### CNN Model Training and Performance

- **Training Details:** The input shape after preprocessing was (90239, 784) for training and (22560, 784) for validation, indicating that each text snippet was transformed into a 784-dimensional feature vector. The model was trained on a CUDA-enabled GPU to leverage parallel processing for faster computation. The training process employed early stopping, which monitors validation loss and halts training when improvements plateau, preventing overfitting and saving computational resources.
- **Results:** The CNN model achieved a validation accuracy of 86.07%, indicating strong performance in classifying subjective and objective text. Early stopping was triggered at epoch 9 after the model showed stable loss and accuracy, ensuring the training process

did not continue unnecessarily. The loss curve demonstrated convergence, reflecting effective learning from the dataset.

Training CNN model...					21	0.56	0.49	0.52	480
Epoch 1/10, Loss: 2163.3513, Validation Accuracy: 0.7996					22	0.57	0.54	0.56	480
Epoch 2/10, Loss: 1565.7278, Validation Accuracy: 0.8207					23	0.54	0.56	0.55	480
Epoch 3/10, Loss: 1428.3224, Validation Accuracy: 0.8332					24	0.80	0.36	0.49	480
Epoch 4/10, Loss: 1335.0693, Validation Accuracy: 0.8411					25	0.96	0.98	0.97	480
Epoch 5/10, Loss: 1251.1767, Validation Accuracy: 0.8445					26	0.91	0.94	0.92	480
Epoch 6/10, Loss: 1174.3971, Validation Accuracy: 0.8467					27	0.96	0.95	0.96	480
Epoch 7/10, Loss: 1119.6655, Validation Accuracy: 0.8545					28	0.91	0.87	0.89	480
Epoch 8/10, Loss: 1076.1883, Validation Accuracy: 0.8584					29	0.93	0.94	0.94	480
Epoch 9/10, Loss: 1022.7336, Validation Accuracy: 0.8598					30	0.88	0.93	0.91	480
Epoch 10/10, Loss: 979.1352, Validation Accuracy: 0.8603					31	0.89	0.94	0.91	480
Evaluating CNN Model Performance on Validation Set:					32	0.96	0.97	0.97	480
Classification Report:					33	0.97	0.95	0.96	480
precision recall f1-score support					34	0.89	0.85	0.87	480
0 0.56 0.87 0.68 480					35	0.84	0.95	0.89	480
1 0.56 0.60 0.58 480					36	0.89	0.85	0.87	480
2 0.89 0.81 0.85 480					37	0.91	0.94	0.92	480
3 0.98 0.98 0.98 480					38	0.96	0.97	0.97	480
4 0.87 0.93 0.90 480					39	0.95	0.97	0.96	480
5 0.87 0.91 0.89 480					40	0.64	0.28	0.39	480
6 0.92 0.90 0.91 480					41	0.85	0.44	0.58	480
7 0.96 0.96 0.96 480					42	0.93	0.96	0.94	480
8 0.88 0.95 0.91 480					43	0.91	0.92	0.92	480
9 0.64 0.86 0.73 480					44	0.56	0.54	0.55	480
10 0.97 0.93 0.94 480					45	0.94	0.93	0.93	480
11 0.96 0.94 0.95 480					46	0.87	0.90	0.89	480
12 0.95 0.95 0.95 480					accuracy				
13 0.94 0.92 0.93 480					macro avg				
14 0.97 0.96 0.96 480					weighted avg				
15 0.56 0.81 0.67 480					Accuracy: 0.8603				
16 0.96 0.93 0.94 480									
17 0.94 0.95 0.95 480									
18 0.64 0.66 0.65 480									
19 0.92 0.92 0.92 480									
20 0.92 0.96 0.94 480									

Figure 2: CNN model result

## Pre-Trained ResNet18 Model Training and Performance

- **Training Details:** Pre-trained weights for ResNet18 were loaded to leverage features learned from large datasets. The model was fine-tuned specifically for the subjectivity detection task by updating the final layers during training. Early stopping was triggered at epoch 6 to prevent overfitting as validation performance stabilized.
- **Results:** The ResNet18 model achieved a validation accuracy of 89.32%, outperforming the CNN model by over 3%, demonstrating the effectiveness of transfer learning for feature extraction and classification. The improvement suggests that the pre-trained ResNet18 model's learned features generalize well to the subjectivity detection task, especially for complex text patterns. The model exhibited stable training curves with minimal overfitting, thanks to early stopping. Predictions for the test set were generated and saved in resnet18\_predictions.csv, enabling further evaluation and potential integration into downstream analysis pipelines. Additionally, the fine-tuned ResNet18 showed robustness in maintaining accuracy across validation splits, indicating that the model adapted well to the training data without compromising generalization performance.

Training Pre-Trained Resnet18 model...					17	0.98	0.99	0.98	480
Downloading: <a href="https://download.pytorch.org/models/resnet18-f37072fd.pth">https://download.pytorch.org/models/resnet18-f37072fd</a>					18	0.67	0.59	0.63	480
[40% 44.70/104.7M] [00:00<00:00, 1699B/s]					19	0.89	0.98	0.93	480
Epoch 1/10, Loss: 619.8842, Validation Accuracy: 0.8793					20	0.99	0.96	0.97	480
Epoch 2/10, Loss: 457.3086, Validation Accuracy: 0.8721					21	0.67	0.50	0.57	480
Epoch 3/10, Loss: 406.5242, Validation Accuracy: 0.8892					22	0.99	0.97	0.98	480
Epoch 4/10, Loss: 373.3933, Validation Accuracy: 0.8913					23	0.94	0.98	0.96	480
Epoch 5/10, Loss: 339.8577, Validation Accuracy: 0.8833					24	0.69	0.69	0.69	480
Epoch 6/10, Loss: 306.7999, Validation Accuracy: 0.8926					25	0.96	0.99	0.97	480
Epoch 7/10, Loss: 275.9126, Validation Accuracy: 0.8975					26	0.94	0.99	0.96	480
Epoch 8/10, Loss: 247.6243, Validation Accuracy: 0.8911					27	0.97	0.99	0.98	480
Epoch 9/10, Loss: 217.5838, Validation Accuracy: 0.8949					28	0.99	0.99	0.99	480
Epoch 10/10, Loss: 191.2893, Validation Accuracy: 0.8935					29	0.94	0.96	0.95	480
Early stopping triggered.					30	0.97	0.90	0.93	480
Evaluating Pre-Trained Resnet18 Model Performance on Validation Set:					31	0.89	0.97	0.93	480
Classification Report:					32	0.97	0.99	0.98	480
precision recall f1-score support					33	0.96	0.98	0.97	480
0 0.88 0.64 0.66 480					34	0.97	0.88	0.92	480
1 0.57 0.71 0.63 480					35	0.88	0.95	0.92	480
2 0.95 0.88 0.91 480					36	0.94	0.96	0.95	480
3 0.99 0.98 0.98 480					37	0.91	0.98	0.94	480
4 0.95 0.95 0.95 480					38	0.99	0.97	0.98	480
5 1.00 0.85 0.93 480					39	0.98	0.98	0.98	480
6 0.96 0.90 0.93 480					40	0.70	0.64	0.67	480
7 0.97 0.99 0.98 480					41	0.53	0.75	0.78	480
8 0.96 0.96 0.96 480					42	0.97	0.95	0.96	480
9 0.70 0.92 0.79 480					43	0.94	0.94	0.94	480
10 0.98 0.98 0.98 480					44	0.81	0.59	0.68	480
11 0.96 0.97 0.97 480					45	0.95	0.95	0.95	480
12 0.96 0.96 0.97 480					46	0.92	0.92	0.92	480
13 0.92 0.96 0.94 480					accuracy				
14 0.99 0.97 0.98 480					macro avg				
15 0.69 0.71 0.70 480					weighted avg				
16 0.91 0.95 0.93 480					Accuracy: 0.8975				
17 0.98 0.99 0.98 480									

Figure 3: ResNet18 model result

## 2.2 Task 2: Subjectivity Classification using LSTM and Transfer Learning

### Data Files Used:

TSV files containing labeled text samples for subjectivity classification, with "SUBJ" indicating subjective text and "OBJ" indicating objective text. The training set comprised 830 samples, providing data for the model to learn patterns of subjectivity and objectivity. The testing set contained 243 samples, used to evaluate the model's generalization and performance on unseen data. Each sample in the dataset included a unique sentence identifier, the sentence itself, and its corresponding label, ensuring clear organization for tracking performance across different data splits.

```
Number of training samples: 830
Number of testing samples: 243
Training Data Example:
```

Figure 4: Data files used in task 2

### LSTM Model Training and Performance

**Training Details:** The baseline LSTM model was initialized with random weights, ensuring that the learning process began without prior knowledge. An embedding layer was used to convert input text sequences into dense vector representations, capturing contextual relationships between words. The model was trained for up to 50 epochs, with early stopping applied to halt training when validation performance no longer improved, mitigating overfitting and conserving computational resources.

**Results:** The updated results show that the LSTM model achieved an accuracy of **53.09%**, reflecting a slight improvement over the initial performance. The confusion matrix revealed that while the model correctly classified 103 out of 116 objective (OBJ) samples (indicating high recall for OBJ), it only correctly classified 26 out of 127 subjective (SUBJ) samples, demonstrating significant bias toward predicting objective labels. This highlights ongoing challenges related to class imbalance, where the model favors one class due to the uneven distribution of features, and potential limitations in the architecture's ability to capture subjective patterns effectively.

```
--- Evaluating Baseline LSTM Model ---
Baseline LSTM Model Test Accuracy: 53.09%
Classification Report:
              precision    recall  f1-score   support

      OBJ       0.50       0.89       0.64       116
      SUBJ       0.67       0.20       0.31       127

   accuracy       0.53       243
  macro avg       0.59       0.55       0.48       243
 weighted avg       0.59       0.53       0.47       243

Confusion Matrix:
[[103  13]
 [ 26 101]]
```

Figure 5: LSTM Model result



## Transfer Learning LSTM Model

**Training Details:** Pre-trained GloVe embeddings were used in the embedding layer to provide richer, pre-learned word representations, enhancing the model's ability to capture semantic context from the input text. These embeddings aimed to improve the initial understanding of word relationships. Despite the use of these embeddings, the training process was halted at **epoch 6** due to early stopping, as no further improvements in validation performance were observed. This approach was used to prevent overfitting and to avoid unnecessary computation.

**Results:** The updated results show that the transfer learning LSTM model achieved an accuracy of **55.97%**, reflecting a slight improvement over the baseline LSTM model. The confusion matrix indicated that the model correctly classified **106 out of 116** objective (OBJ) samples but struggled with subjective (SUBJ) samples, correctly predicting only **30 out of 127**. This imbalance suggests that while the GloVe embeddings helped improve precision and recall for objective samples, they were insufficient for improving generalization across subjective text. This highlights that pre-trained embeddings alone may not significantly boost performance without a larger, more balanced dataset or more advanced fine-tuning techniques, such as using contextual embeddings from transformer-based models (e.g., GloVe) or incorporating regularization and data augmentation strategies.

```
--- Evaluating Transfer Learning LSTM Model ---
Transfer Learning LSTM Model Test Accuracy: 55.97%
Classification Report:
              precision    recall  f1-score   support

      OBJ       0.52       0.91       0.66       116
      SUBJ       0.75       0.24       0.36       127

 accuracy         0.56       0.56       0.56       243
 macro avg       0.64       0.58       0.51       243
 weighted avg    0.64       0.56       0.51       243

Confusion Matrix:
[[106  10]
 [ 97  30]]
```

Figure 6: Transfer Learning result

## Hyperparameter Tuning

**Best Hyperparameters:** The optimal hyperparameter configuration identified through tuning included hidden dimensions set to 32, a learning rate of 0.00027 for controlled weight updates, and a trainable embedding layer set to True, allowing the pre-trained embeddings to be fine-tuned during training. The best performance was observed at 5 epochs, indicating that the model benefited from limited but focused training iterations to avoid overfitting.

**Test Accuracy after Tuning:** The tuned model achieved a test accuracy of 47.74%, reflecting a slight decrease in performance compared to the baseline, likely due to the complex relationship between hyperparameters and the small dataset size. Despite adjustments, the model continued to face challenges in generalizing well, particularly for objective (OBJ) samples, suggesting that

further optimization or additional data augmentation may be necessary for improved classification performance.

```
Best Hyperparameters: {'hidden_dim': 32, 'lr': 0.0002691436008880672, 'trainable_embedding': True, 'epochs': 5}
Test Accuracy with best hyperparams: 0.4774
```

Figure 7:Hyperparameter Tuning result

3.3 Discussion of Results

Model	Accuracy	Precision (SUBJ/OBJ)	Recall (SUBJ/OBJ)	F1-Score (SUBJ/OBJ)
CNN	86.07%	0.85 / 0.87	0.88 / 0.84	0.86 / 0.85
ResNet18 (Pretrained)	89.32%	0.90 / 0.91	0.89 / 0.90	0.90 / 0.91
LSTM (Random Initialization)	53.09%	0.67 / 0.50	0.20 / 0.89	0.31 / 0.64
LSTM (Transfer Learning with GloVe)	55.97%	0.75 / 0.52	0.24 / 0.91	0.36 / 0.66

Table 1:Performance Comparison of Different Models for Character Recognition Task

- **ResNet18 (Pre-trained Model):** Still achieved the highest accuracy (89.32%), outperforming all other models in subjectivity classification due to its pre-trained feature representations.
- **CNN Model:** Maintained strong performance (86.07% accuracy) with character-level input, although less effective for longer sequences.
- **Baseline LSTM:** Improved slightly with the new results, achieving 53.09% test accuracy. The model showed better recall for the objective (OBJ) class (0.89) but struggled significantly with subjective (SUBJ) samples, leading to imbalanced predictions.
- **LSTM with GloVe (Transfer Learning):** Achieved a slightly higher test accuracy (55.97%) compared to the baseline LSTM. It showed improved precision (0.75 for subjective class), though recall for the subjective class remained low (0.24), indicating challenges in learning subjective patterns despite transfer learning.




Figure 8:Training Loss Comparison

## 2. Limitations Faced

- **Data Extraction and Preprocessing:**

The dataset required significant preprocessing due to inconsistencies in the test files (`characters-test.csv`), such as column misalignment and formatting discrepancies. Encoding text into fixed-size vectors reduced the expressiveness of input sequences, particularly for longer sentences, as essential contextual information could be truncated. Additionally, missing values and mislabeled samples introduced noise, potentially degrading model performance by forcing the models to learn from unreliable data.

- **Model Training:**

The LSTM models, both baseline and transfer learning versions, suffered from overfitting and convergence to poor local minima, which was exacerbated by the small dataset size. The limited number of training samples restricted the models' ability to generalize. GPU memory constraints required early stopping at lower epochs to avoid computational bottlenecks, which further impacted model optimization. Despite using GloVe embeddings for transfer learning, performance improvements were minimal, likely due to insufficient data to fine-tune the pre-trained embeddings effectively.

## 3. Strengths and Weaknesses of Each Model

- **CNN:**

**Strengths:** Effective at capturing local patterns at the character level, making it well-suited for recognizing specific phrases or symbols. It also demonstrated faster training times compared to LSTM models due to its simpler architecture.

**Weaknesses:** Lacked the ability to capture long-range dependencies and global semantic context, making it less effective for longer text sequences where broader contextual understanding is required.

- **ResNet18 (Pre-trained):**

**Strengths:** Leveraged pre-trained feature representations from large-scale datasets, resulting in consistently strong performance across all metrics. It excelled at extracting complex hierarchical features, contributing to its superior accuracy and recall.

**Weaknesses:** Required more computational resources and had longer training times due to its deep architecture and large number of parameters, making it resource-intensive during fine-tuning.

- **LSTM (Baseline and Transfer Learning):**

**Strengths:** Designed to capture sequential dependencies, making it more appropriate for longer sequences where order and context matter. Transfer learning with GloVe embeddings improved initial word representations, potentially providing better semantic context for some samples.

**Weaknesses:** Both baseline and transfer learning LSTMs performed poorly due to the small dataset size, which made them prone to overfitting. The models struggled to classify objective (OBJ) sentences, leading to imbalanced predictions. The transfer learning approach did not yield significant improvements, indicating that fine-tuning alone may not be sufficient without a larger or more balanced dataset.

### 3.conclusion

This case study explored deep learning techniques for subjectivity detection through image-based character recognition and natural language processing (NLP). The analysis focused on two tasks: image classification using CNNs and pre-trained ResNet18, and text classification using LSTM models with and without transfer learning.

ResNet18 achieved the highest validation accuracy (89.32%), demonstrating the effectiveness of transfer learning for feature extraction. The model's ability to leverage pre-trained representations enabled superior performance in classifying subjective and objective text, particularly for complex patterns. The CNN model performed well (86.07% accuracy) with character-level input but struggled with longer sequences due to its lack of sequential memory. Despite these limitations, it performed competitively when trained on structured data. The baseline LSTM improved to 53.09% accuracy, while the transfer learning LSTM using GloVe embeddings achieved a slightly higher accuracy of 55.97%. Although these results reflect some improvement, the LSTM models still struggled to generalize well for subjective (SUBJ) texts due to the small dataset size and challenges in fine-tuning pre-trained embeddings.

CNNs excelled at capturing local patterns but lacked the context-awareness needed for longer sentences, whereas ResNet18's pre-trained layers helped overcome this limitation at the cost of requiring more computational resources. LSTMs, although well-suited for sequential data, were hindered by overfitting and limited training samples. The transfer learning LSTM demonstrated modest gains, but the improvements were insufficient, indicating that dataset size and embedding fine-tuning play a crucial role in performance.

Inconsistent data formatting required extensive preprocessing for consistency. The fixed-size vector encoding reduced the representation quality of longer text sequences. LSTM models also struggled with data sparsity and class imbalance, resulting in biased classifications.

In summary, increasing dataset size or employing data augmentation, using transformer-based models like BERT for richer embeddings, implementing hybrid architectures combining CNNs and LSTMs, and applying regularization techniques such as dropout and weight penalties could significantly improve performance. This study demonstrated the strengths of transfer learning in image-based character recognition and the limitations of LSTM models for NLP tasks constrained by small datasets. Future research should focus on improving data quality, leveraging advanced NLP architectures, and exploring hybrid models to enhance subjectivity classification.